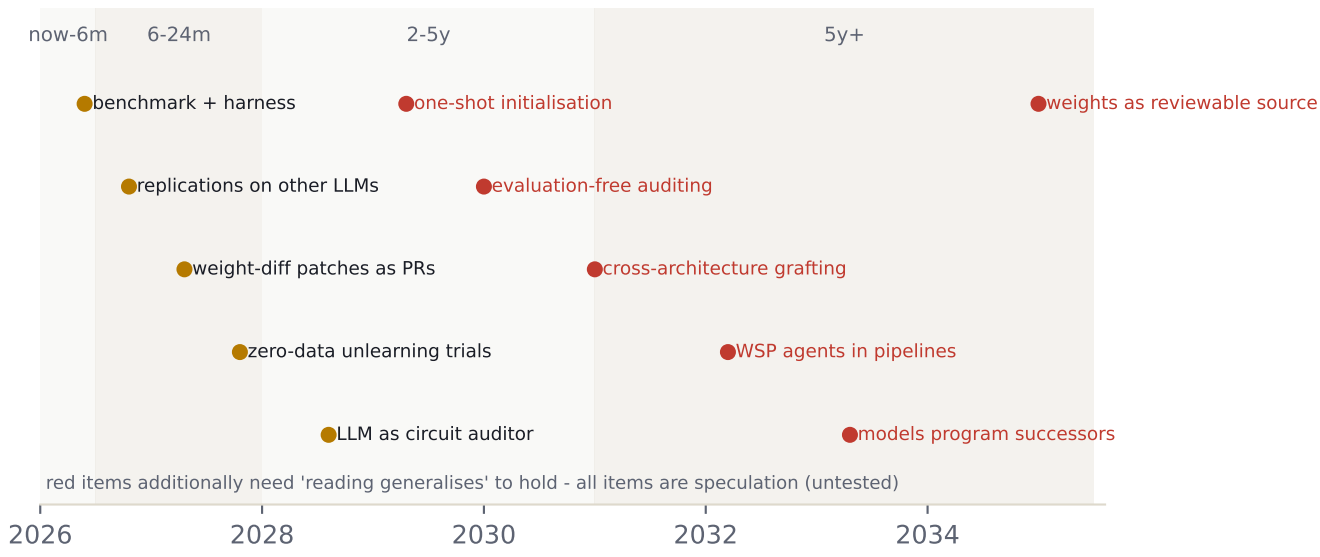
now-6m
6-24m
2-5y
5y+
benchmark + harness
one-shot initialisation
weights as reviewable source
replications on other LLMs
evaluation-free auditing
weight-diff patches as PRs
cross-architecture grafting
zero-data unlearning trials
WSP agents in pipelines
LLM as circuit auditor
models program successors
red items additionally need 'reading generalises' to hold - all items are speculation (untested)
2026
2028
2030
2032
2034